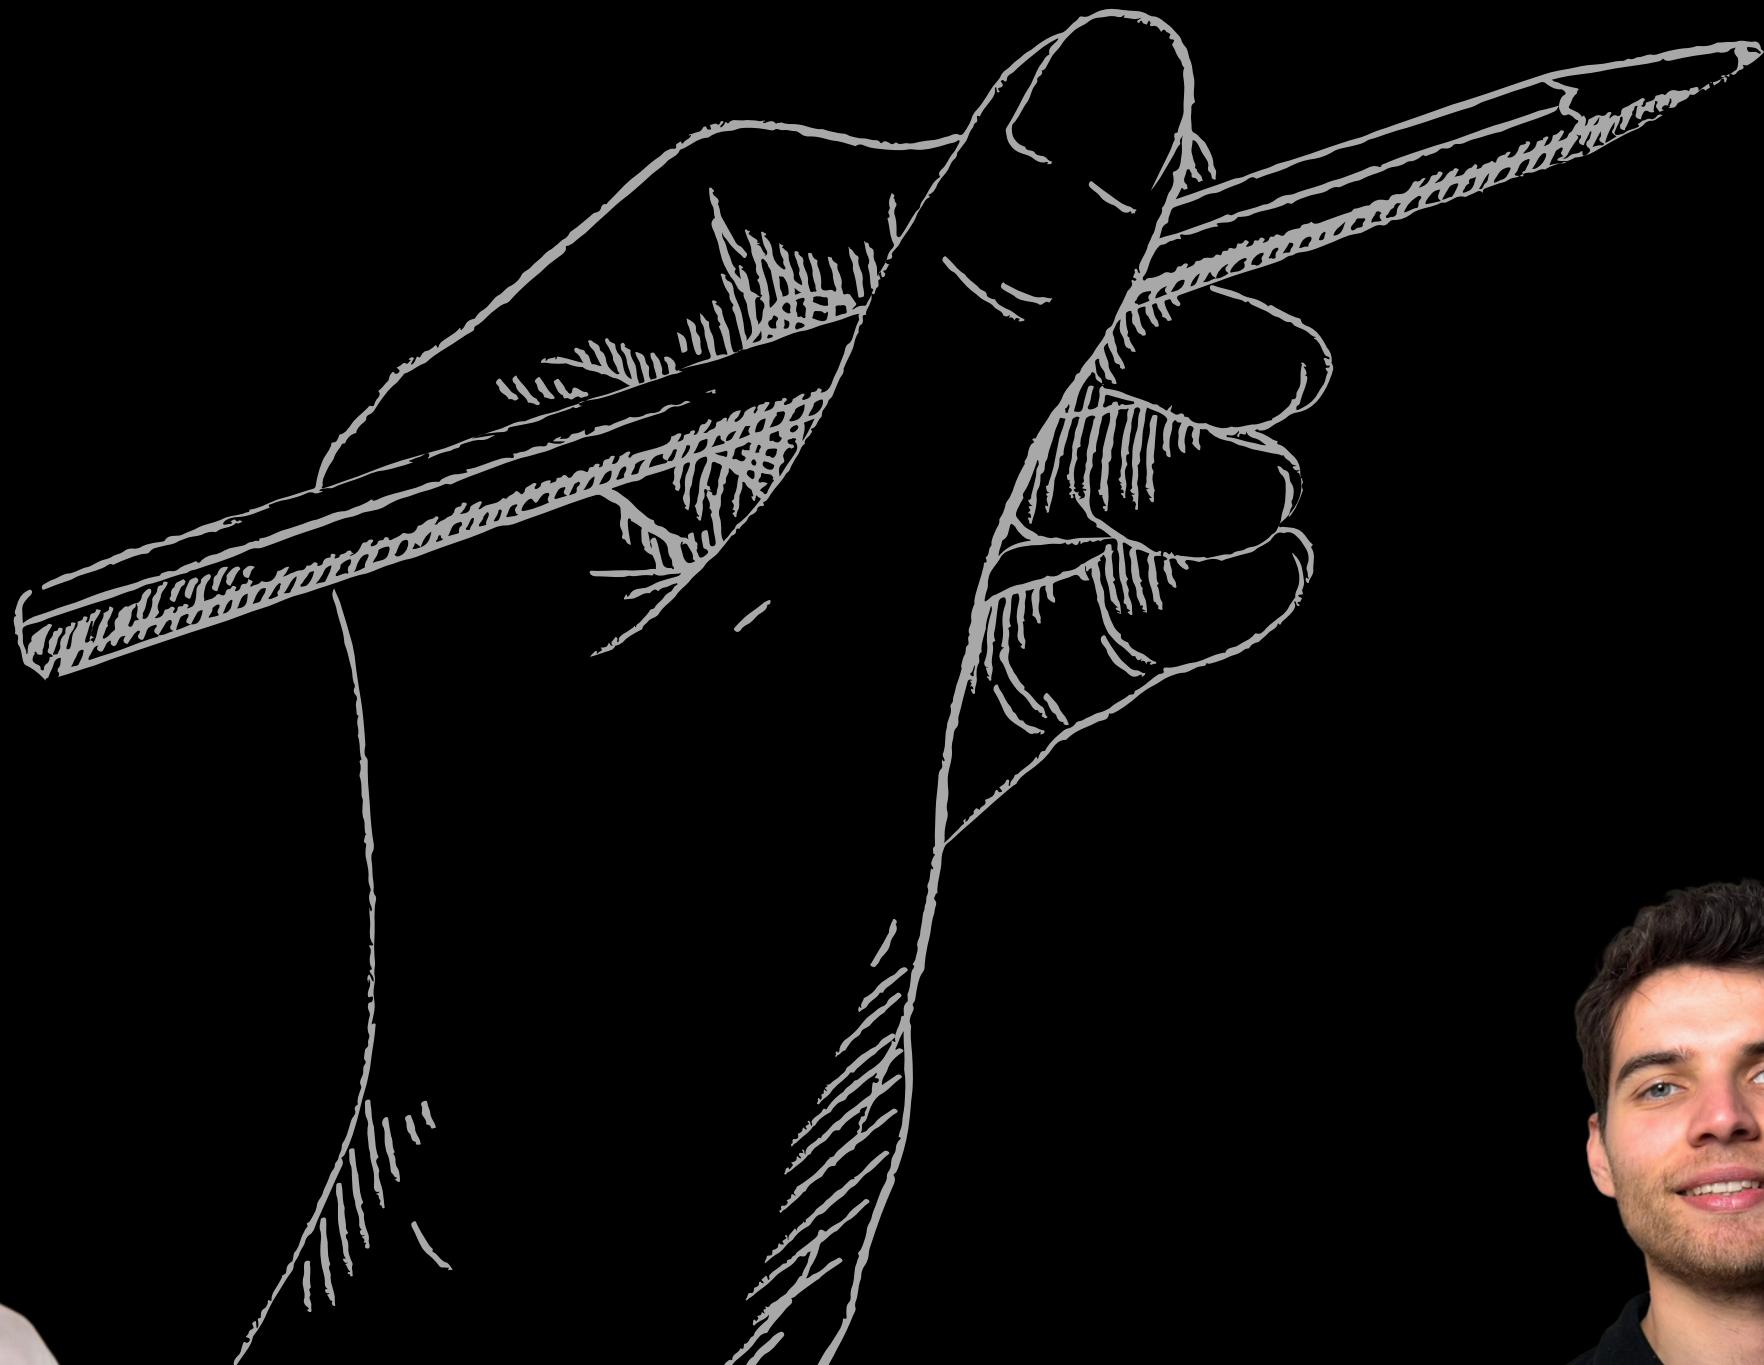


ML END2END PIPELINE

TIPS FROM REAL-
WORLD PROJECTS



INTRO / PURPOSE

An end-to-end machine learning project aims to take a model from idea to real-world use, including data pipelines, deployment, monitoring, and updates.

A static project in a notebook usually stops at training and testing the model, without worrying about how it runs in production or keeps working over time.

The end-to-end approach makes the work useful outside the notebook.



TABLE OF CONTENT

1. Setting env (Python env, libraries, Git/GitHub...)
2. Data ingestion
3. Data validation
4. **EDA** (Exploratory data analytics)
5. Feature Engineering
6. Model Training
7. Evaluation
8. Modular System Logic
9. Python scripts (load, train, ...)
10. FastAPI endpoint
11. Docker image
12. CI/CD pipeline
13. Deploy
14. Monitor
15. Retraining Loop
16. Common Mistakes



1/ SET ENVIRONMENT

A

PROJECT STRUCTURE

It's important to set a project structure before starting your project. Some components appear in most ML project, make sure to include them.

```
mkdir -p \
  data/raw data/processed data/external \
  notebooks \
  src/{data,features,models,utils} \
  app \
  configs \
  scripts \
  tests \
  .github/workflows \
  docker \
  great_expectations \
  mlruns \
  artifacts
```

```
✓ TELCO CUSTOMER CHURN MLE
  > .github
  > .gradio
  > .venv
  > artifacts
  ✓ configs
  > data
  > docker
  > great_expectations
  > mlruns
  > notebooks
  > scripts
  > src
  ✓ tests
  ⚡ .dockerignore
  ⚡ .gitignore
  ⏪ CLAUDE.md
  ⚡ dockerfile
  ⓘ README.md
  Ⓜ requirements.txt
```

Run this command in project's terminal to create all the structure for you.



Why? Because 'most' ML projects have a similar structure so no need to keep doing the same thing over and over ➤

1/ SET ENVIRONMENT

B

INSTALL LIBRARIES AND CREATE VIRTUAL ENVIRONMENT

Run this command in project's terminal to create **requirements.txt**. Include all needed libraries.

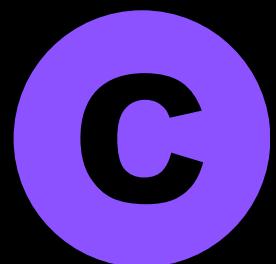
```
cat > requirements.txt << 'EOF'  
pandas  
numpy  
scikit-learn  
mlflow  
fastapi  
unicorn  
pydantic  
python-dotenv  
joblib  
great-expectations  
pytest  
EOF
```

```
python3.11 -m venv .venv  
source .venv/bin/activate  
pip install --upgrade pip  
uv pip install -r requirements.txt
```

Run this command in project's terminal to create **virtual environment** and install libraries using '**uv**'. **uv** in this case speeds up the installation.

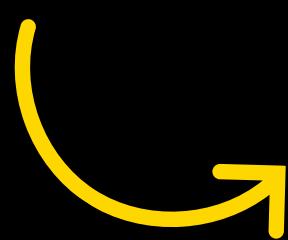
Why? Each project needs a dedicated environment to avoid dependency conflicts.

1/ SET ENVIRONMENT



SET GIT & GITHUB

Go to GitHub.com, sign in to your account, and create a new repository.



Run these commands in project's terminal to initialise GIT, add a Readme file, commit your first change, and link your project to GitHub



Create a new repository Preview S

Repositories contain a project's files and version
Required fields are marked with an asterisk (*).

1 General

Owner * Repository name *

 anesriad / My_AWESOME_REPO

My_AWESOME_REPO is a public repository

Great repository names are short and memorable

Description

Let me show you my awesomeness :)

35 / 350 characters

...or create a new repository on the command line

```
echo "# My_AWESOME_REPO" >> README.md
git init
git add README.md
git commit -m "first commit"
git branch -M main
git remote add origin https://github.com/anesriad/My_AWESOME_REPO.git
git push -u origin main
```

Why? Each project needs a dedicated environment to avoid dependency conflicts.



2/ DATA INGESTION

Before you analyze or model anything, you need to load your data reliably and reproducibly.

Where data comes from

- CSVs, SQL/NoSQL databases, APIs, cloud storage

Light preprocessing at this stage

- Rename columns, parse dates, drop duplicates
- Apply consistent schemas or dtypes early

Reproducibility

- Use loading scripts, not manual steps
- Version input data if possible

```
# load_data.py

import pandas as pd
from dotenv import load_dotenv
import os

load_dotenv() # Load .env variables

DATA_PATH = os.getenv("DATA_PATH", "data/raw/customers.csv")

def load_data(path=DATA_PATH):
    df = pd.read_csv(path)
    df.columns = df.columns.str.lower().str.strip()
    df.drop_duplicates(inplace=True)
    return df
```

Why? Because broken inputs = broken pipeline. Clean, consistent ingestion is the foundation of every ML project. ➤

3/ DATA VALIDATION

You want to validate incoming data before engineering features or training.

WHAT TO CHECK

- Are column names/types correct?
- Are there missing or out-of-range values?
- Do values match expected formats?

HOW TO VALIDATE

- Use libraries like pydantic, Great Expectations, or pandas assertions.
- Run validation early so errors don't silently break your model later.

```
# Showcase of Great Expectations for data validation

import great_expectations as ge

# convert pandas DataFrame to a Great Expectations dataset
ge_dataframe: ge.dataset.PandasDataset = ge.from_pandas(df)

# check that contract_type contains only expected categories
contract_type_expectation_result: dict = ge_dataframe.expect_column_values_to_be_in_set(
    column="contract_type",
    value_set=["month-to-month", "one year", "two year"]
)
print("Contract type categories valid:", contract_type_expectation_result["success"])

# check that monthly_charges are within a reasonable range
monthly_charges_expectation_result: dict = ge_dataframe.expect_column_values_to_be_between(
    column="monthly_charges",
    min_value=10,
    max_value=200
)
print("Monthly charges within range:", monthly_charges_expectation_result["success"])
```

Best practice
If validation fails, fail fast and stop the pipeline.

Why? Bad data silently breaks your pipeline. Validation catches issues early, before they corrupt training or hit production.

4/ EXPLORATORY DATA ANALYTICS

A

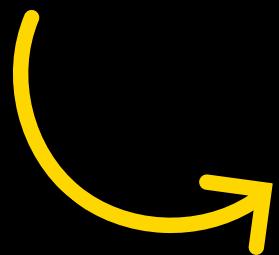
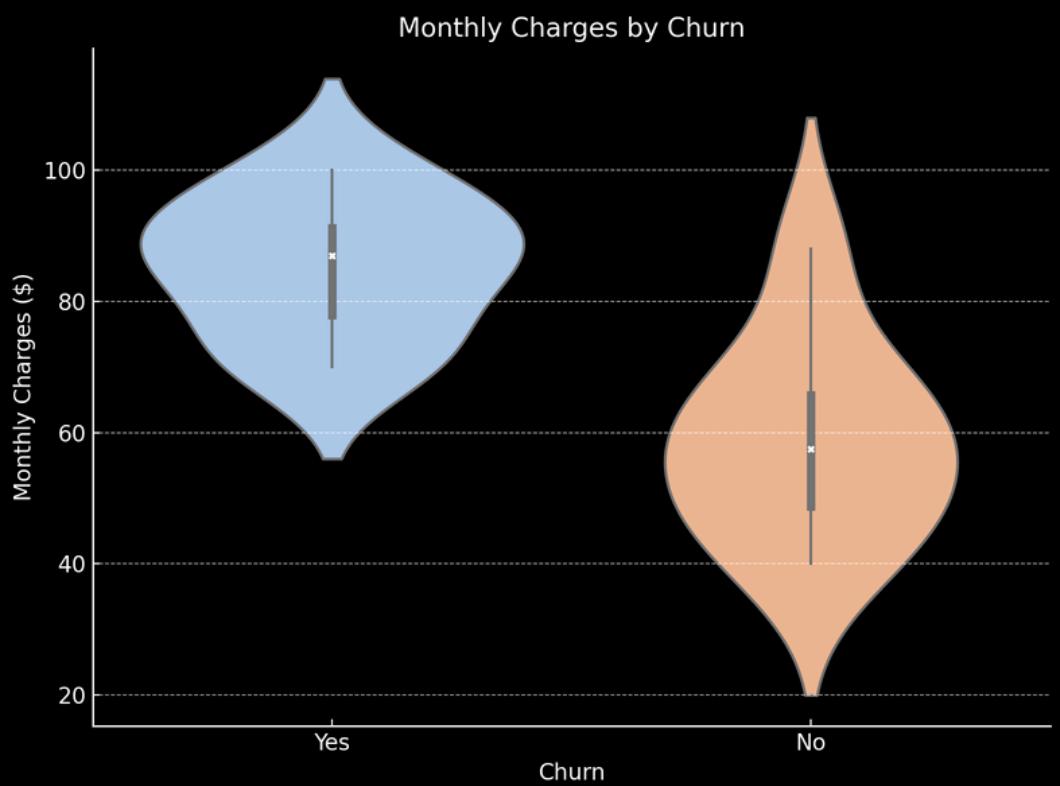
WHY EDA?

UNDERSTAND • SPOT PROBLEMS • FORM HYPOTHESES

EDA isn't just "looking at data."

It's how you:

- Understand distributions and relationships
- Spot missing values and outliers
- Catch broken logic in your dataset
- Shape the path for modeling



TYPES OF EDA

UNIVARIATE

→ Understand distributions and relationships

BIVARIATE

→ Scatter plots to spot relationships

MULTIVARIATE

→ Correlation heatmaps for variable interactions

TIME SERIES

→ Line plots to detect trends and seasonality

Start simple, then go deeper as patterns emerge.

Why? EDA helps you understand your data, spot issues early, and decide what to do next.



4/ EXPLORATORY DATA ANALYTICS

B

STRUCTURED EDA IN PRACTICE

3 steps to make your EDA faster, cleaner, and actually useful

1/ Dataset Overview

Understand your data before diving in (before your first plot).

- **Rows & columns:** Do you have 10K rows or 10M? It changes everything.
- **Data types:** Explicitly verify int, float, object, bool, datetime.
- **Missingness & Duplicates:** What % is missing? Are duplicates valid?
- **Target variable distribution:** Look for imbalance early.

2/ Feature Deep-Dive

Now zoom in. Focus on distributions and transformations.

- **Numerical:** Outliers? Skewed? Is normalization needed?
- **Categorical:** Too many levels? Are rare classes meaningful or noise?
- **Zero-heavy fields:** Are they real zeros, or just unused? (e.g. income = 0)
- **Leakage risk:** Any feature too good to be true?
- **Correlations:** Find redundancies before modeling.

3/ Quality Traps & Broken Logic

EDA is debugging for your data, not just describing it.

- Inconsistent records: Conflicting info across rows?
- Invalid categories: Typos, inconsistent casing, mixed encodings.
- Linked nulls: Are two features always missing together?
- Suspicious values: Negative ages? Future timestamps?

In churn prediction, EDA might reveal patterns like how “last login” or “support tickets opened” vary between retained vs churned users

Why? A structured flow makes your EDA faster, repeatable, and easier to debug.



5/ FEATURE ENGINEERING

Transform raw data into signals your model can learn from

TYPICAL STEPS

- Handle missing values
- Encode categoricals
- Normalize numeric features
- Create new features from domain logic
- Aggregate or transform time-based features
- Drop redundant or leaking features

FEATURE ENGINEERING

Numerical Features

- Normalize values
- Bin into ranges (e.g., low/med/high usage)
- Create ratios or differences between columns

Categorical Features

- Encode with one-hot or label encoding
- Group rare categories
- Map to risk levels if applicable

Date Features

- Extract durations (e.g., account age)
- Calculate time since events
- Flag recent activity

Text Features

- Count keyword mentions
- Detect sentiment or intent
- Encode using embeddings or TF-IDF

Interaction Features

- Combine fields to reveal patterns
- Multiply related columns (e.g., usage × cost)
- Use domain logic to create new variables



Avoid Data Leakage

Only use info available at prediction time.
No future payments, no labels, no outcomes.

Why? Smart features make patterns easier to learn and often have the biggest impact on model performance.

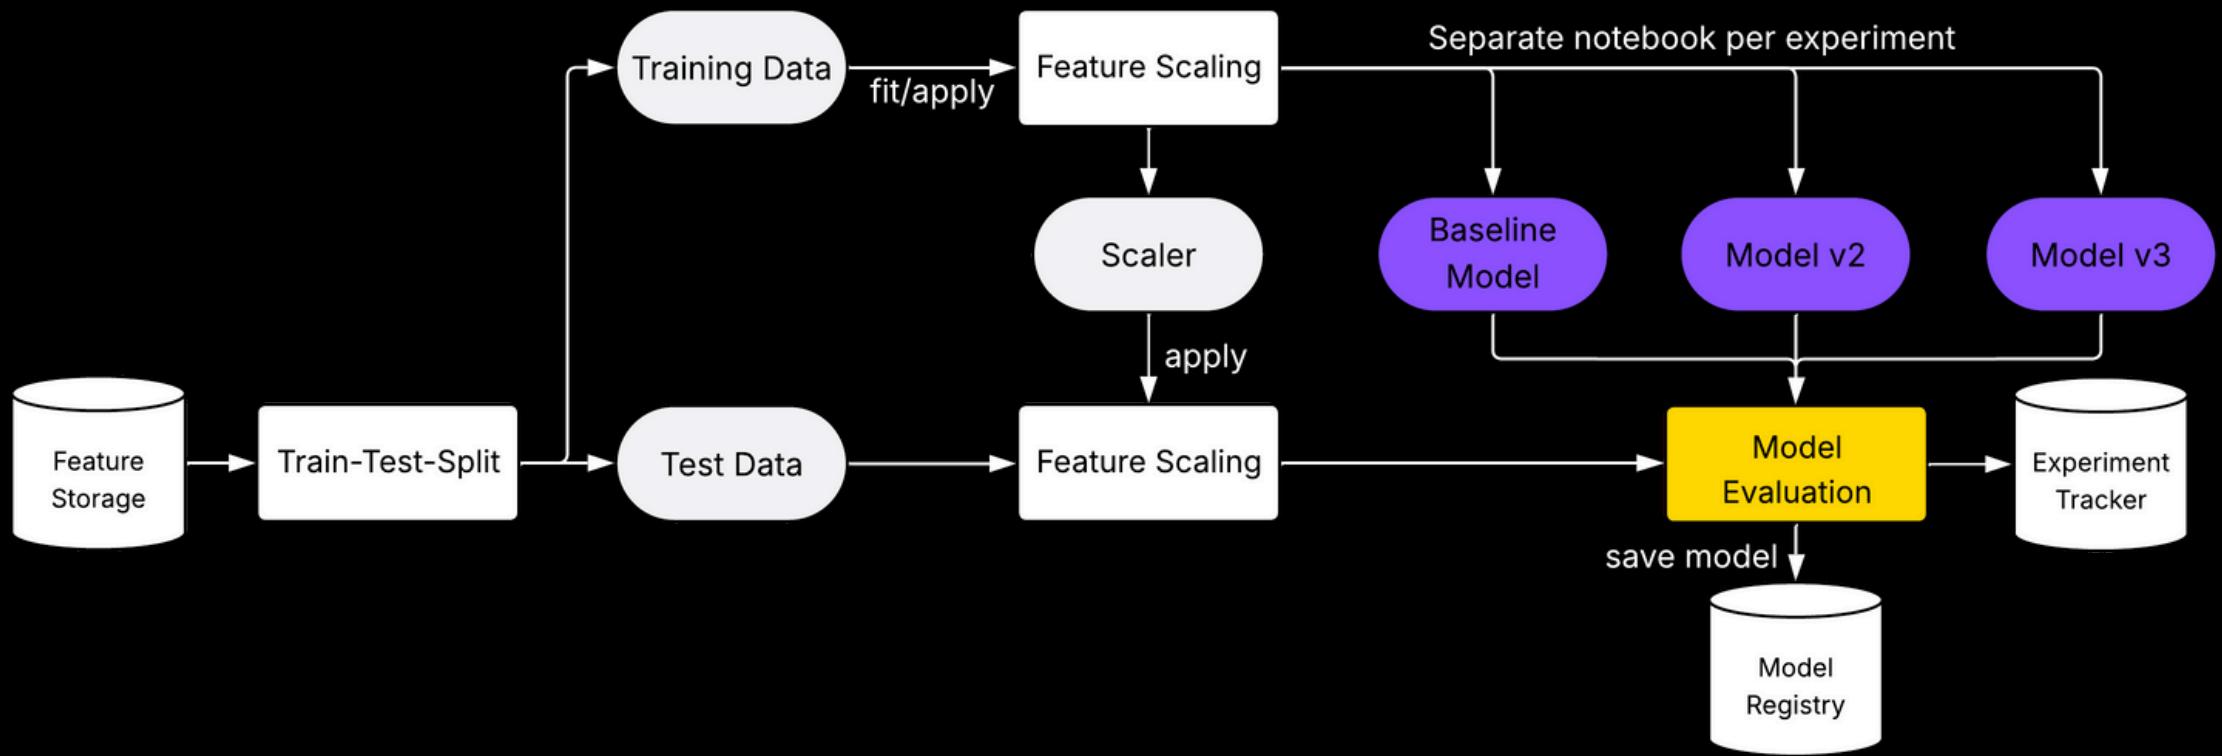


6/ MODEL TRAINING

Model training means teaching your algorithm to generalize from past data to unseen data.

How it works:

- Split data into training + test sets
- Fit scalers on training data only
- Train a baseline model (e.g., logistic regression)
- Try new architectures or hyperparams
- Log results to track what works
- Evaluate and version your best model



Example (Churn):

You might start with logistic regression, then try XGBoost or LightGBM as iterations. Track every result in your experiment tracker (e.g., MLflow)

Why? Training helps your model learn patterns, tracking helps you learn which patterns worked best.

7/ EVALUATION

MEASURE WHAT MATTERS

Evaluation tells you how well your model performs and whether it's ready for production.

What to evaluate:

- Accuracy or F1 → for classification (like churn prediction)
- MAE / RMSE → for regression
- AUC → when class imbalance matters
- Confusion matrix → for error analysis
- Calibration → if probabilities are used in decisions

Tips

- Always compare on the same test set
- Track metrics across all experiments
- Use plots (ROC, PR curve, residuals) to dig deeper
- Go beyond averages: inspect where your model fails

Example (Churn):

Your model might have 85% accuracy, but most of that comes from predicting "no churn". F1, F2 and confusion matrix will give a clearer picture.



What to log:

- Metrics
- Test dataset version
- Model version
- Dataset Hash
- Hyperparameters used

Why? Evaluation reveals how your model performs, and where it breaks.



7/ MODEL REGISTRY & EXPERIMENT TRACKING

TRACKING EXPERIMENTS

helps you avoid “notebook chaos” and compare what’s actually working. A model registry makes your deployment traceable and versioned, like **Git for models**.



WHAT TO TRACK:

- Parameters (e.g., learning rate, max_depth)
- Metrics (e.g., F1, ROC-AUC)
- Artifacts (e.g., trained model files, plots)
- Code version & dataset snapshot

HOW TO USE IT:

- Use **MLflow** to log training runs with `mlflow.log_params()` and `mlflow.log_metrics()`
- Register the best-performing model in the Model Registry
- Deploy only models that passed evaluation + review

Tip: Tag each model with metadata like `production_ready=True`, so it's clear what's deployable.

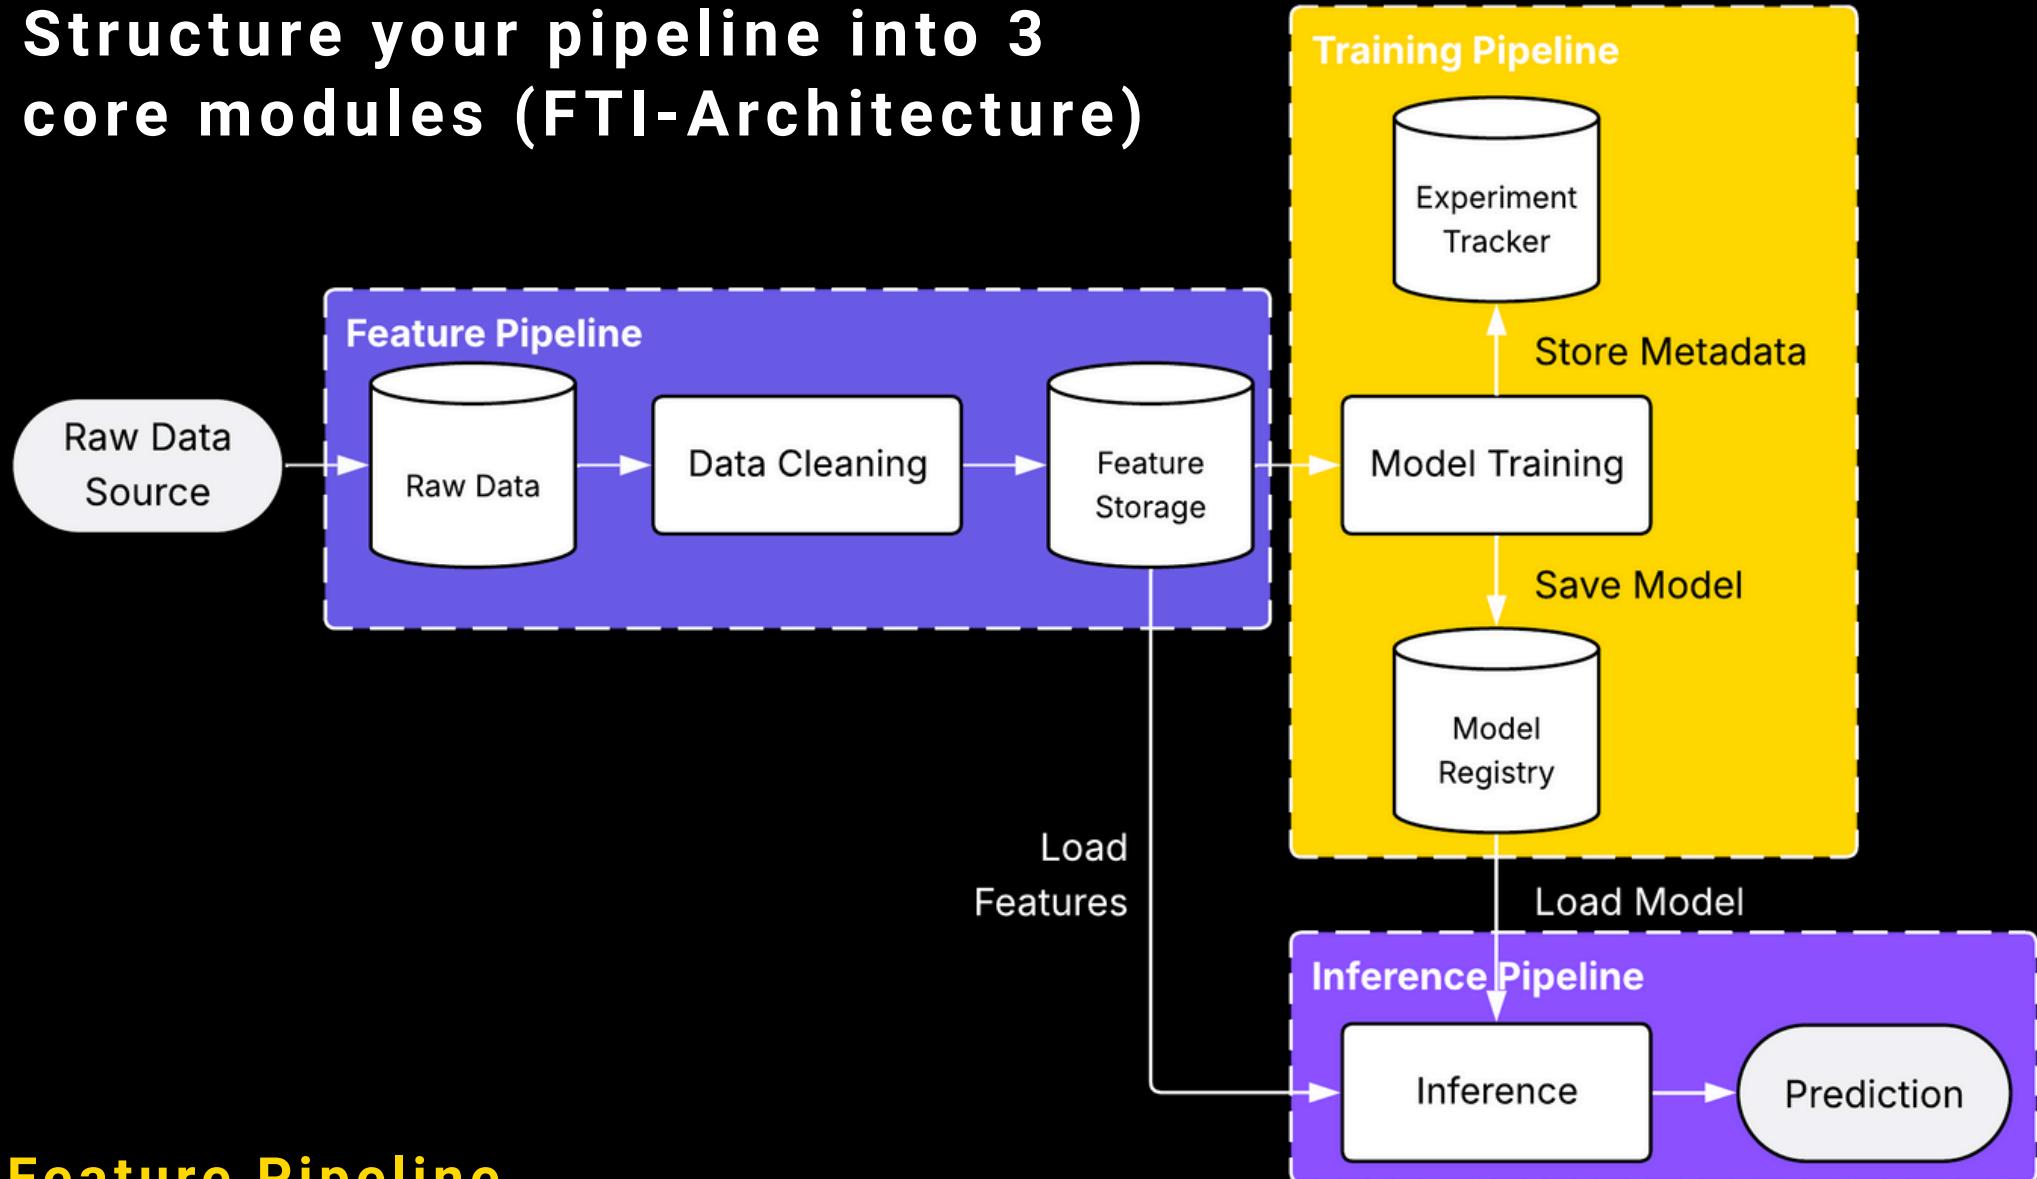
Why? Evaluation reveals how your model performs, and where it breaks.



8/ MODULARIZE LOGIC

Tangled pipelines break fast.

Structure your pipeline into 3 core modules (FTI-Architecture)



Feature Pipeline

- Prepares reusable features for both training and inference.
- Ingestion: Handle raw source data & schemas
- Preprocessing: Clean, normalize, encode
- Feature Engineering: Transformations split by data type or domain
- Output: Logged to a feature store (optional)

Training Pipeline

- Focuses on model development, reproducibility, and evaluation.
- Training: Model selection, hyperparams, config mgmt
- Evaluation: Metrics, diagnostics, drift checks
- Output: Saved model & metadata

Inference Pipeline

- Serves predictions in production environments.
- Feature Retrieval: Pull latest features or recompute
- Serving: Wrap model with APIs, format outputs
- Output: Scored results, downstream integrations

Why? Modular pipelines make your ML system easier to scale, test, and maintain.

9/ MODULARIZE LOGIC

```
✓ notebooks
  └─ 01_data_description.ipynb
  └─ 02_exploratory_data_analysis.ipynb
  └─ 03_data_cleaning.ipynb
  └─ 04_data_encoding.ipynb
  └─ 05_time_series_analysis.ipynb
  └─ 06_train_test_split.ipynb
  └─ 07_naive_forecast.ipynb
  └─ 08_lgbm_forecast.ipynb
```

Core Logic

*separate notebooks for each processing step and experiment



Keep IO, logic, and config separate. Treat each stage of the architecture as a replaceable unit, not a one-off script.

```
✓ src
  ✓ config
    ! base.yml
    ! inference.yml
    ! training.yml
  ✓ evaluation
    ✎ evaluate.py
    ✎ metrics.py
    ✎ plots.py
  ✓ features
    ✎ feature_store.py
  ✓ inference
    ✎ api_wrapper.py
    ✎ postprocessing.py
    ✎ predict.py
  ✓ ingestion
    ✎ data_loader.py
    ✎ source_config.py
  ✓ preprocessing
    ✎ clean.py
    ✎ encode.py
    ✎ normalize.py
  ✓ training
    ✎ model_config.py
    ✎ model_factory.py
    ✎ scheduler.py
    ✎ train.py
  ✓ utils
    ✎ io.py
    ✎ logging.py
```

Shared config defaults
Service-time behavior
Model + training setup

Run eval pipeline
Custom scoring logic
Model diagnostics

Store and load features

Serve model via API
Tidy outputs for users
Run model inference

Pull and load source data
Raw data setup

Fix raw data issues
Encode categoricals
Scale numeric features

Set hyperparameters & options
Build models from config
Trigger + manage training
Run model training

Read/write data and models
Central logging logic

Why? Modular code lets you reuse parts, test in isolation, and swap logic without breaking the whole.



10/ STARTER WORKFLOWS

BUILD REAL SKILL

TRY THESE MINI WORKFOWS:

1. Notebook → Scripted Pipeline

Start with a notebook → extract code into preprocess.py, train.py, predict.py → run with main.py or CLI
Mental bridge from one-off notebook to repeatable projects

2. Train → Save → Predict

Train model → save with joblib.dump() → load later with joblib.load() in new script

Introduces serialization and persistance, a must-know before deployment

3. FastAPI Wrapper (Single File)

Load a model in FastAPI → define /predict endpoint → accept JSON → return predictions

Shows how a model becomes a real service

4. GitHub + CI Test

Push to GitHub → GitHub Actions runs test: e.g. “python predict.py works”

Easy intro to CI/CD; teaches that your code should always be runnable

5. CLI For Feature Engineering

Wrap preprocessing into CLI script:

- “python build_features.py --input raw.csv --output features.csv”

Promotes modularity and real-world thinking (users, parameters, file paths).

Why? These workflows train you to move beyond notebooks and think like a system builder

Why? Pushes learners to structure and parameterize logic – a big skill jump.

10/ FASTAPI ENDPOINT

FastAPI acts as the **serving layer**, it turns our ML model into an API that can **receive data** and **return predictions**.

It's necessary because without it, the model would just sit in Python files, and no external app, script, or user could interact with it.

```
# filename: main.py
from fastapi import FastAPI
from pydantic import BaseModel
import joblib
import numpy as np

# create FastAPI app
app = FastAPI()

# define request format
class InputData(BaseModel):
    features: list[float]

# load a trained model (replace with your model file)
# For example, a scikit-learn model saved with joblib
model = joblib.load("model.joblib")

@app.get("/")
def home():
    return {"message": "ML model is ready!"}

@app.post("/predict")
def predict(data: InputData):
    """
    Take input features -> run through model -> return prediction
    """

    X = np.array(data.features).reshape(1, -1) # convert list to 2D array
    prediction = model.predict(X)[0]
    return {"prediction": float(prediction)}
```

The diagram illustrates the flow of code execution. It starts with the import statements at the top, followed by the creation of the FastAPI app. Below that, the request format is defined. Then, the code loads a trained model. Finally, two endpoints are set up: a GET endpoint at the root URL (/) and a POST endpoint at /predict. Purple arrows point from the right margin to the corresponding code lines, indicating the flow of execution from right to left.

imports the FastAPI framework.

creates the web app object.

sets up a GET endpoint at the root URL (/).

sets up a POST endpoint at /predict.

Why? *FastAPI makes it easy and fast to turn your ML model into a web API that others can call for predictions.*

11/ CONTAINERIZE (DOCKER)

Docker lets us package the whole project (code, dependencies, environment) into one container so it runs the same everywhere. The goal is for the project to run on any machine, not just yours.

create 'dockerfile' in your repo

git dockerfile > ...

```
You, 2 days ago | 1 author (You)
1 # 1- Use a slim Python image
2 FROM python:3.11-slim
3
4 # 2- Set working directory inside the container
5 WORKDIR /app
6
7 # 3- Copy requirements file and install dependencies
8 COPY requirements.txt .
9 RUN pip install --no-cache-dir -r requirements.txt
10
11 # 4- Copy the rest of the app code
12 COPY . .
13
14 # 5- Expose FastAPI's default port
15 EXPOSE 8000
16
17 # 6- Run the app with unicorn
18 CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8000"]
```

6

19

1

2

3

4

5

To start the container:

* churn-api is the name of the container to start

```
docker build -t churn-api .
docker run -p 8000:8000 churn-api
```

Why? Docker makes ML projects portable and reproducible by packaging code, dependencies, and models into one container that runs anywhere.

12/ CI/CD - GITHUB ACTIONS

A CI/CD pipeline **automates testing, building, and deployment** so code changes are checked and delivered quickly without manual steps.

GitHub Actions lets you set up this pipeline directly in your repo, running workflows (like tests or deployments) whenever you push code. Add this, when your model is stable and you are deploying regularly.

```
name: Build and Push to Docker Hub  
  
on:  
  push:  
    branches: [main]  
  
jobs:  
  build-and-push:  
    runs-on: ubuntu-latest  
  
    steps:  
      - name: Checkout code  
        uses: actions/checkout@v4  
  
      - name: Set up Docker Buildx  
        uses: docker/setup-buildx-action@v3  
  
      - name: Log in to Docker Hub  
        uses: docker/login-action@v3  
        with:  
          username: ${{ secrets.DOCKERHUB_USERNAME }}  
          password: ${{ secrets.DOCKERHUB_TOKEN }}  
  
      - name: Build and push Docker image  
        uses: docker/build-push-action@v5  
        with:  
          context: .  
          file: ./dockerfile  
          push: true  
          tags: anasriad/maroc-fastapi:latest
```

in your project create
→ `.github/workflows/ci.yml`

Only trigger on main branch pushes

modern Docker builder that supports multi-platform builds

Set API token username & password on Docker Hub.
Save those secrets in GitHub → Secrets & variables

Build & push docker image to docker hub

Why? It lets you automatically build and publish your container to Docker Hub so anyone (or any server) can pull and run the latest version of your ML app easily.

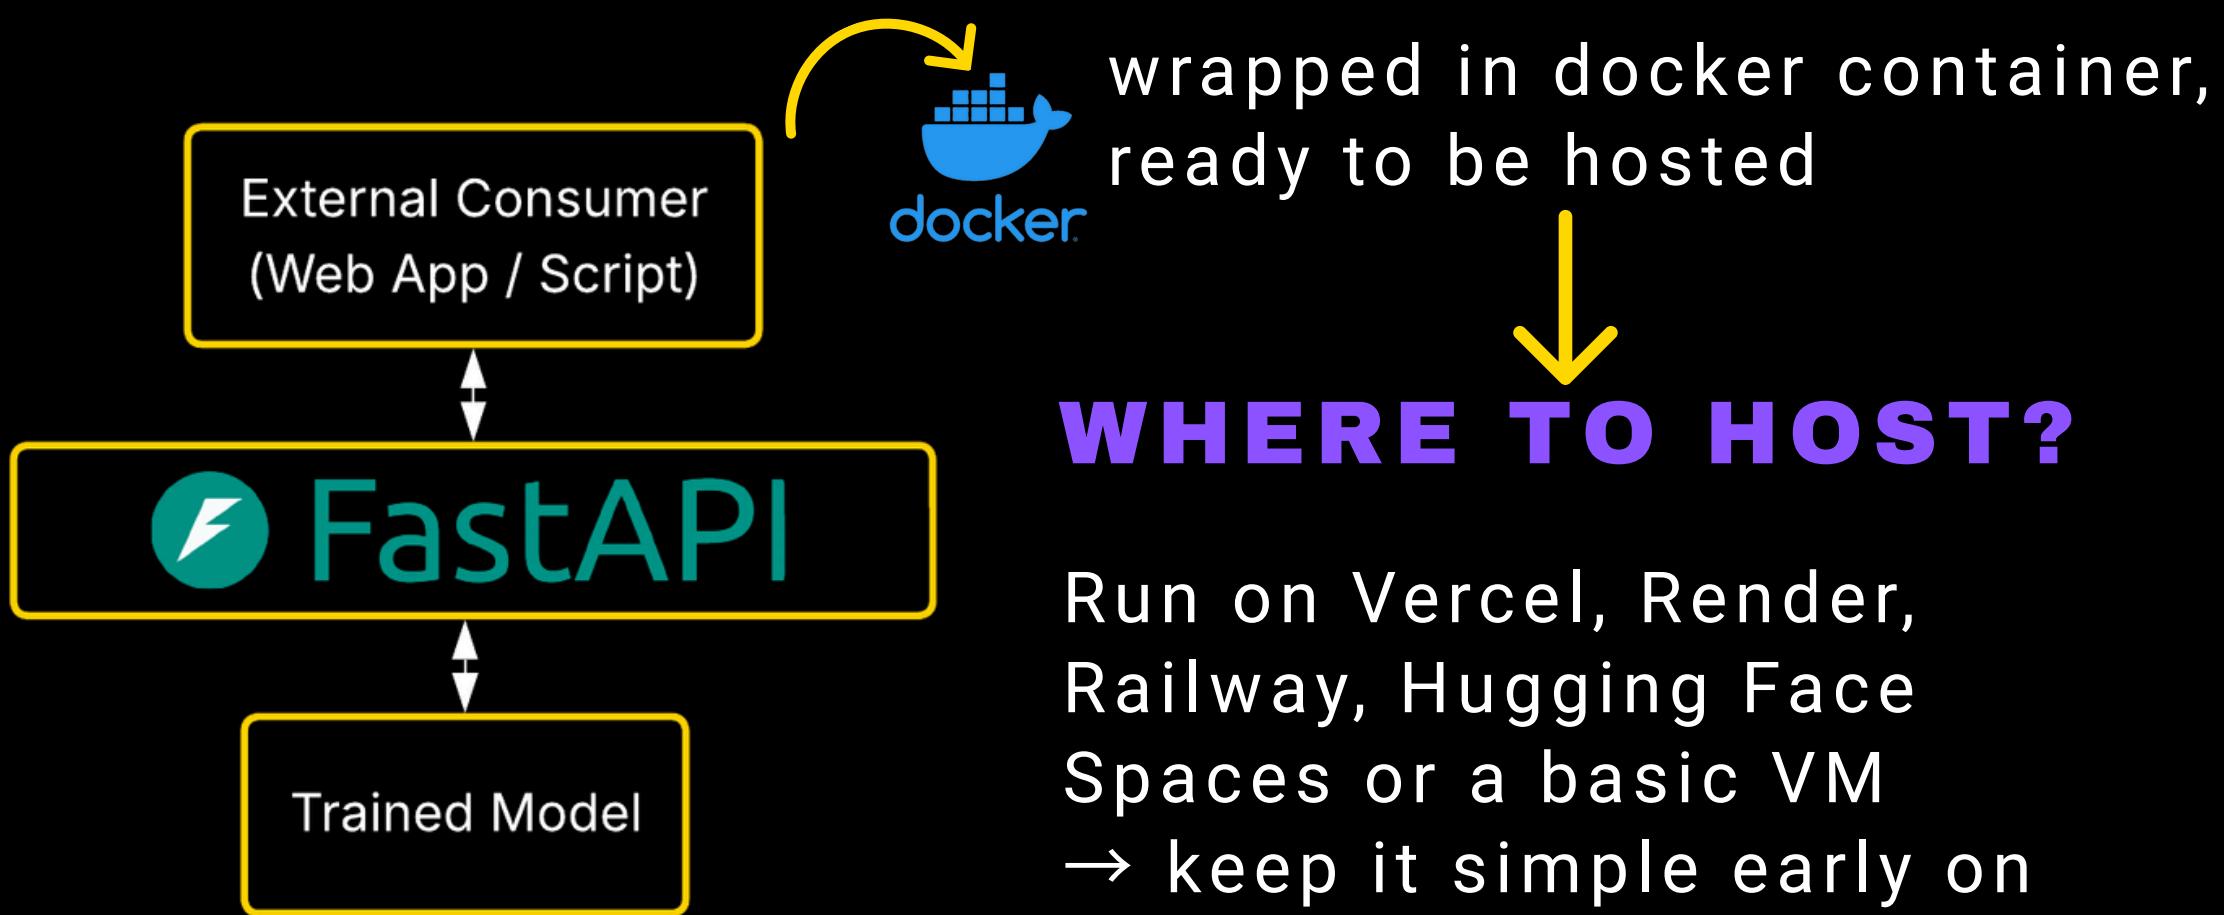
13/ DEPLOY

After training, your model is just a file. Deployment turns it into a real product that delivers value.

Monitoring ensures it keeps delivering value.

- Use FastAPI to serve predictions via a REST API
- Use .env files to separate dev/prod configs
- Docker Compose to run app + model together
- MLflow to version and track deployed models
- Run dummy requests post-deploy to catch silent failures

Return predictions as JSON, CSV, or database updates depending on consumer



why? Deployment isn't just "putting a model online", it's making it stable, fast, and easy to use for others.

14/ MONITOR

- **Data Drift:** Input distributions change over time. Log stats like mean/std for key features
- **Prediction quality:** Track metrics (e.g. accuracy, F1) and output distribution
- **Service health:** Uptime, latency, errors
- **Observability:** Use logging tools and set up alerts (e.g., with Prometheus + Grafana)
- **Traceability:** Store requests/responses for debugging and include inputs, model version and timestamps



Why? Monitoring helps you fix silently failing systems before users notice.

15/ RETRAINING LOOP

Why retrain?

→ Model quality drops over time (new users, new patterns, outdated logic)

What can trigger it?

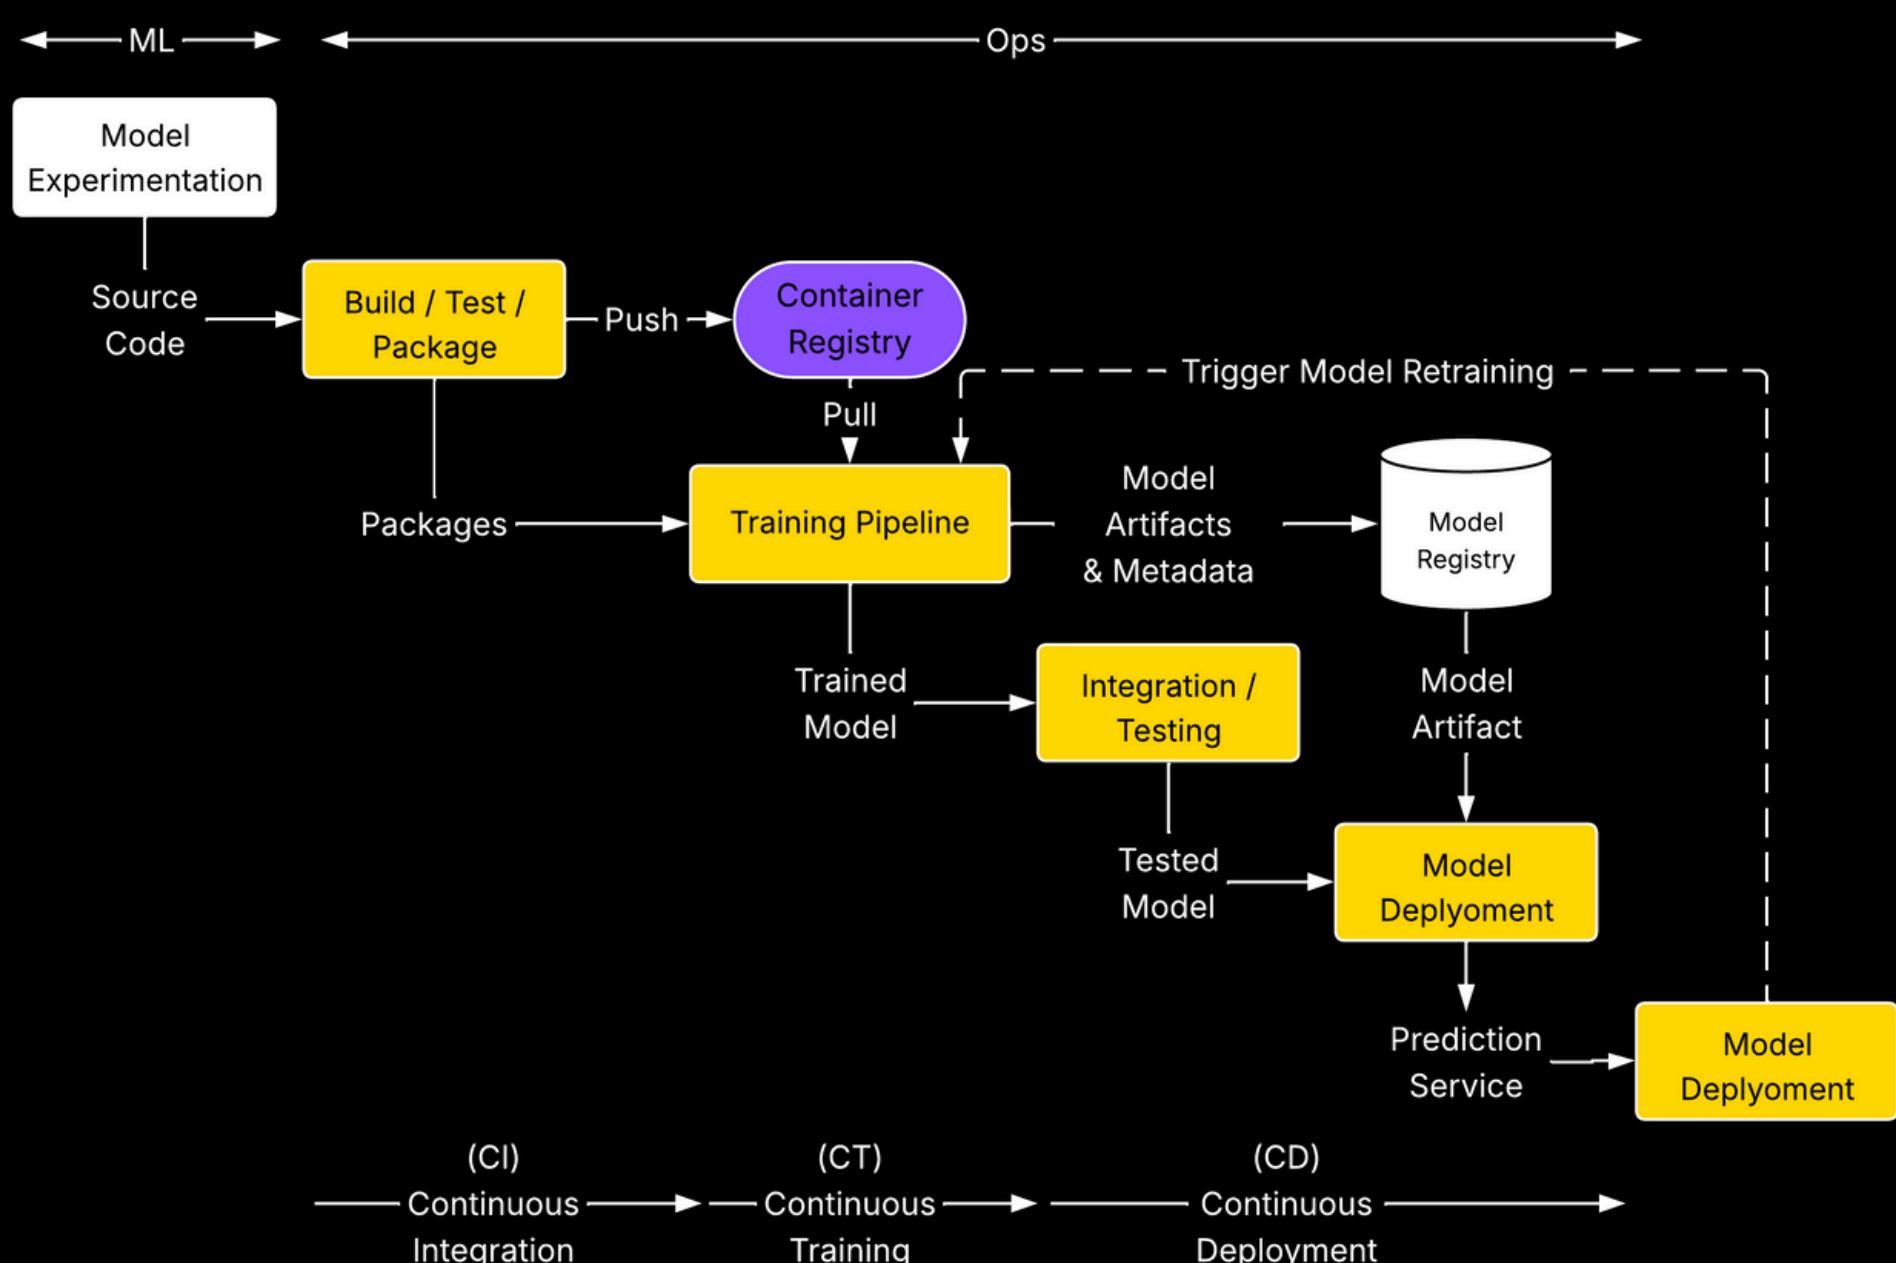
→ Data drift, accuracy drop, new data, schema changes

How to automate it:

→ Cron job, GitHub Actions, or an Airflow DAG

What happens during retraining?

→ Pull fresh data → rerun pipeline → train → evaluate → register new model



Why? ML systems degrade over time, a retraining loop keeps them relevant and accurate.

16/ COMMON MISTAKES

Most ML mistakes are small, but costly.

They often hide in setup, not in the modeling itself.

Data Leakage:

You used future data in training without realizing it.

No Model Versioning

You deployed a model, but no one knows which one.

Only Trained on Training Data

No test set, no generalization.

Manual Retraining

You reran the notebook and silently broke something.

Silent Failures

The API returns results, but they're wrong, and no one knows.

Drift Ignored

Your input data changed months ago, but your model didn't.

No Monitoring

Something goes wrong, but there are no alerts.

Too Many Experiments, No Tracking

You tried 20 things, but forgot which one worked best.



**WHICH STAGE YOU
FOUND THE MOST
INTERESTING?**

COMMENT BELOW



**FOLLOW ANAS AND MARKUS
FOR MORE CONTENT ON DATA
SCIENCE, AND ML
ENGINEERING.**

